

1.6 Consistency Algorithm

Consistency Algorithm

Exercise 1.

- Clearly (2) \Rightarrow (3) since it is always true that

$$\pi_{\bar{x}_i}(S_1 \bowtie \dots \bowtie S_n) \subseteq S_i \bowtie S_j \subseteq S_i$$

- If (1) then an homomorphism h witnessing $q \rightarrow D$ gives suitable $S_i = \{h(\bar{u}_i)\}$, showing (1) \Rightarrow (2).
- Suppose (2), and pick $\bar{u} \in S_1 \bowtie \dots \bowtie S_n$, which is non empty because the S_i are non empty. We can deduce from \bar{u} a map $h_{\bar{u}}$: it maps variable x to the single element in the projection $\pi_{(x)}(\bar{u})$. Then clearly $h_{\bar{u}}(\bar{x}_i) \in S_i \subset q_i(D)$ and hence $h_{\bar{u}}$ is a homomorphism from q_i to D for every i . We get finally that h is a homomorphism witnessing $q \rightarrow D$, which shows (2) \Rightarrow (1).
- Suppose finally (3) and let us show (2). Choose $s_i \in S_i$, and deduce homomorphisms h_i witnessing $q_i \rightarrow D$. We can also pick a join tree T rooted at X_i . Our goal here is to show that $s_i \in \pi_{\bar{x}_i}(S_1 \bowtie \dots \bowtie S_n)$, and we will show that we can choose suitable h_j for all j that produce consistency.
 - First, if $X_k = X_j$, and h_k is already chosen, we can choose $h_j = h_k$ and it works, since $h_k(\bar{x}_i)$ is consistent with some tuple in S_j .
 - Thus it is enough to choose h_j only for nodes in T . We claim that it is enough to choose h_j so that it is consistent with the homomorphism associated to the parent of X_j . The proof is basically the same as in the proof of proposition 19.2 in the book.

Hence we have shown for all i that $S_i \subset \pi_{\bar{x}_i}(S_1 \bowtie \dots \bowtie S_n)$, and the inverse inclusion being always true, this concludes.

Other solution, prove by induction on the join tree structure that (local) implies (global)

Exercise 2. Start with $S_i = q_i(D)$, and iterate the operation of iterating over all i and performing a semi-join of S_i on all other S_j , successively.

This works because if $A \subset B$ and $C \subset D$ and $A = A \bowtie C$, then $A \subset B \bowtie D$.

Exercise 3. $q = R(a, b)R(b, c)R(c, a)$

R[1]	R[2]
1	2
2	1

Yannakaki's Algorithm

Exercise 4. Sort both and do a joined scan over $q(D)$ and $q'(D)$, which computes in $\mathcal{O}(x \log(x) + y \log(y))$. Then to compute

$$q_S = \text{Answer}(\bar{y}) : -R_{j_1}(\bar{u}_{j_1}) \dots$$

we iteratively compute

$$\begin{aligned} & R_{j_1}(\bar{u}_{j_1})(D) \\ & R_{j_1}(\bar{u}_{j_1})(D) \bowtie R_{j_2}(\bar{u}_{j_2})(D) \\ & \vdots \end{aligned}$$

Exercise 5.

- ★ First we show that the first pass can be computed in the desired complexity. We can first compute all $q_s(D)$. A single computation takes $\mathcal{O}(\|q_s\| \cdot \|D\| \cdot \log(\|D\|))$ so the sum of the computations for all s is $\mathcal{O}(\|q\| \cdot \|D\| \cdot \log(\|D\|))$. Remains to compute the

$$Q_s(D) = \bigcap_{1 \leq i \leq p} q_s(D) \times Q_{s_i}(D)$$

We recall that $\|q_s(D)\| \leq \|D\|$ and $\|Q_{s_i}(D)\| \leq \|q_{s_i}(D)\| \leq \|D\|$, implying that $Q_s(D)$ can be computed in $\mathcal{O}(p \cdot \|D\| \log \|D\|)$. The sum of those computations for all s asks to control the size of each p , which sum to the number of nodes in the join tree – that is $\mathcal{O}(\|q\|)$. In the end, all Q_s can be computed in $\mathcal{O}(\|q\| \cdot \|D\| \cdot \log(\|D\|))$.

- ★ Note that if \bar{x} is contained in a single node of the join tree, call it t , which can be detected in $\mathcal{O}(\|q\|)$, then we can root the join tree at t and execute only the first pass of the algorithm. Then, the computed $Q_t(D)$ is already $q(D)$ (up to a projection onto \bar{x} which can be done in $\mathcal{O}(\|Q_t(D)\|) = \mathcal{O}(\|D\|)$). Hence in that case, the total running time of the algorithm is $\mathcal{O}(\|q\| \cdot \|D\| \cdot \log(\|D\|))$.
- ★ Next we show that the second pass can be computed in the desired complexity. This is easy : $\|A_s\| \leq \|Q_s(D)\| \leq \|D\|$, where we get that A_s can be computed in $\mathcal{O}(\|D\| \log(\|D\|))$. In the end, the whole second pass can be computed in $\mathcal{O}(\|q\| \cdot \|D\| \cdot \log(\|D\|))$.
- ★ Next we show that in the second pass of the algorithm, the sets A_s that are computed are the answers on D to the queries

$$A_s(\bar{y}_i) : -R_1(\bar{u}_1) \dots R_n(\bar{u}_n)$$

that is, we prove theorem 20.3 of the book. Note that for clarity we write A_s for the sets that are computed by the algorithm and A_s for the queries defined above. We show the result by induction on the structure of the tree. Clearly for the root, this is true. Now suppose $A_s = A_s(D)$ and let s' be a child of s . If $\bar{a} \in A_{s'}(D)$ then clearly $\bar{a} \in A_{s'}$ (indeed $A_{s'} \subseteq Q_{s'}$). The difficulty lies in showing that $A_{s'} \subseteq A_{s'}(D)$.

Let $\bar{a} \in A_{s'} = Q_{s'}(D) \times A_s(D)$, from which we get a homomorphism $h_{s'}$ witnessing $(A_{Q_{s'}}, \bar{y}_{s'}) \rightarrow (D, \bar{a})$. By definition, there is a tuple $\bar{b} \in A_s(D)$ that is consistent with \bar{a} . This in turns means there is a homomorphism h witnessing $(A_q, \bar{y}_s) \rightarrow (D, \bar{b})$. Finally, define h' by completing $h_{s'}$ with h . We claim that h' witnesses $(A_q, \bar{y}_{s'}) \rightarrow (D, \bar{a})$. Clearly we only need to show that it witnesses $A_q \rightarrow D$. Consider some atom $R_i(\bar{y})$ in q . We write $\bar{y} = \bar{y}_1 \cup \bar{y}_2 \cup \bar{y}_3$ where

- \bar{y}_1 are the variables of \bar{y} that are not in $Q_{s'}$
- \bar{y}_2 are the variables of \bar{y} that are shared between s' and s
- \bar{y}_3 are the variables of \bar{y} that are in $Q_{s'}$ but not in s' . Note that then those variables do not appear anywhere else but in the subtree rooted in s' , because of the join tree property.

Because of the join tree property, \bar{y}_1 and \bar{y}_3 cannot be simultaneously non empty. Moreover, h and $h_{s'}$ are equal over \bar{y}_2 . This shows that h' is equal to either h or $h_{s'}$ over \bar{y} , whence we get that $R_i(h'(\bar{y})) \in D$.

To summarize, we have shown that $\bar{a} \in A_{s'}(D)$, which proves that $A_{s'} = A_{s'}(D)$.

- ★ Finally, we must bound the complexity of the third pass. There is some ambiguity as to what it means to project X onto a dimension that is not present in X , i.e. we compute $\pi_{s \cup \bar{x}}(O_s(D) \bowtie O_{s_j}(D))$ but we are not certain that all variables of \bar{x} appear in $s \cup s_j$. In that case, I understood that the variables of \bar{x} that are not represented are ignored. If s is a node of the join tree, let \bar{Y}_s be the variables' subset of $\bar{x} \cup s$ that appear in the subtree rooted in s . Define the queries

$$O_s(\bar{Y}_s) : -R_1(\bar{u}_1) \dots R_n(\bar{u}_n)$$

We first note that the O_s that the algorithm computes are the $O_s(D)$. Indeed, thanks to the join tree property, $O_s = \pi_{s \cup \bar{x}}(A_s(D) \bowtie_{s_j \text{ child of } s} O_{s_j}(D))$, i.e. we can defer the projection to the very end. Thus $O_s = A_s(D) \bowtie \pi_{s \cup \bar{x}}(\bigbowtie_{s_j \text{ child of } s} O_{s_j}(D))$. In particular $O_s \subseteq A_s(D) \bowtie \pi_{\bar{x}}(\bigbowtie_{s_j \text{ child of } s} O_{s_j}(D))$ whence we get that $\|O_s(D)\| = \mathcal{O}(\|D\| \cdot \|q(D)\|)$.

We also notice that in the course of its computation, O_s only ever increases, and thus its size is always bounded by $\|O_s(D)\|$.

And here I am stuck...

Acyclicity of the Core

Exercise 6. $q(y_1, \dots, y_n) = \exists x \exists z R(x, y_1) \dots R(x, y_n) R(z, y_1) \dots R(z, y_n)$

Or a 2 cycle and any $2n$ cycle

Exercise 7. If R_i in q is mapped to R'_j in q' , then define S_i as such : for every $\bar{a} \in S'_j$, deduce what \bar{x}_i is mapped to through h , call it $\eta_i(\bar{a}) = (\eta_{x_{i_1}}(\bar{a}), \dots, \eta_{x_{i_p}}(\bar{a}))$ and add that to S_i , where h witnesses $q \rightarrow q'$. If R_i is only mapped to constants, add that one tuple of constants to S_i .

Then take $\bar{u} \in S_i$, choose some $j \neq i$, its associated $\bar{a} \in S'_i$, and $\bar{b} \in S'_j$, consistent with \bar{a} . Finally, note $\bar{v} \in S_j$ associated with \bar{b} . We want to show that \bar{u} and \bar{v} are consistent.

If $x_{i_k} = x_{j_l}$ is mapped to a variable, it works because \bar{a} and \bar{b} are consistent, if it is mapped to a constant then it works directly. In short, it works.

Conjunctive Queries with Equational Atoms